


Bayesian Optimization for Field Nutrient Mapping [Concept] [Draft]



1. Data Processing & Feature Engineering

The "All-In" Feature Set

Instead of feeding only mid-points, we feed a rich feature vector into the GPyTorch model.

- **Target (y):** Nutrient values (Nitrogen, etc.).
- **Coordinates (x_s):** Latitude, Longitude (normalized).
- **Time (x_t):** Days since a reference date (e.g., start of season).
- **Covariates (x_c):** * **CatBoost/XGBoost Embeddings:** Extract the last hidden layer or leaf indices. These act as "latent features" that describe the soil's chemistry beyond just spatial proximity.
 - **Topography:** Slope, Aspect, and Elevation (DEM).
 - **Satellite Priors:** NDVI/NDRE values from Sentinel-2.

Handling Noise & Multi-Sampling

- **Individual Points:** Each raw sensor measurement is kept.
 - **Weighted Likelihood:** For a point with n samples, we set the **Fixed Noise Variance** to $\sigma^2_{\text{device}} / n$. GPyTorch allows you to pass an `observed_noise` vector directly into the Gaussian Likelihood.
-

2. GPyTorch Architecture: The Multitask Spatio-Temporal Kernel

We use a **Multitask Kernel** combined with a **Product Kernel** to handle space, time, and relationships between nutrients.

The Kernel Recipe

1. **Spatial Core:** `Matern-3/2` kernel. It handles the "jumps" you mentioned better than a smooth RBF.
2. **Temporal Core:** `Matern-1/2` (for jagged time changes) multiplied by a `PeriodicKernel` (to capture seasonality).
3. **Cross-Nutrient Logic:** `MultitaskKernel` using **Intrinsic Coregionalization (ICM)**. This allows the model to learn that if Nitrogen spikes, Phosphorus usually follows.
4. **Decay/Weights:** To make recent points more important, we multiply the spatial kernel by a **Time-Decay Kernel**. Points from 2 years ago will have a high spatial correlation but a low "Temporal-Trust" score.

$$K_{\text{total}} = (K_{\text{space}} \otimes K_{\text{time}}) \times K_{\text{nutrients}}$$

3. Implementation Blueprint

Phase 1: Global Training (Offline)

- **Objective:** Learn "Universal" hyperparameters (e.g., the average noise of the device, the typical correlation length of Nitrogen).
- **Method:** Use **Stochastic Variational Gaussian Processes (SVGP)** on your entire global database. Use **Inducing Points** (\$M \approx 500-1000\$) to represent the millions of historical data points.
- **Tools:** GPyTorch on a GPU instance.

Phase 2: Local Update (Real-time Serving)

- **Objective:** Update the map for a specific field in milliseconds.
 - **Method:** 1. Load the Global Model.
2. Use `get_fantasy_model()` or a **Local Exact GP** for the specific field's points (since \$N < 10,000\$).
3. Treat Lab Data as **Constraints**: Set their noise variance to nearly zero (\$10^{-6}\$).
-

4. Serving & System Design

- **Backend:** Python (FastAPI) running the GPyTorch inference engine.

- **Storage:** * **PostgreSQL/PostGIS** for coordinates.
 - **S3** for the pre-trained Global Model "Weights" (Hyperparameters).
 - **Workflow:**
 - a. User takes a sample → App sends Raw Data + Sensor Mean.
 - b. Backend queries Satellite/Topography data.
 - c. GPyTorch performs the Bayesian Update.
 - d. Returns: **Mean_Map** (for display) and **Uncertainty_Map** (for active sampling).
-

5. Estimates & ROI

Category	Detail
Execution Time	PoC: 3 weeks. Production: 4 months.
Prediction Speed	Sub-500ms per field update.
Hardware	Standard Cloud Instance (CPU for inference, GPU for global training).
Reliability	Self-Correcting: High noise results in high uncertainty, preventing bad decisions.
User Experience	"The Guide": The UI tells the user exactly where to poke next to save time.

6. What you haven't considered (The "Secret" Layer)

Heteroscedastic Noise Learning:

Don't just assume the device is noisy. Let the model learn *why* it's noisy. By feeding in **soil_moisture** and **temp** as covariates, the model can learn: *"When the soil is >40% moisture, the sensor noise triples."* The map will automatically reflect this by increasing the "blur" in wet areas, protecting the farmer from bad data.

1. The "Ultimate" Hybrid Model: Hierarchical GPR

Instead of choosing between raw points and means, use a **Hierarchical Gaussian Process**.

- **How it works:** You don't "flatten" the data before the app. You feed the model the raw sensor time-series (or the embeddings from your XGBoost) as *likelihoods*.
 - **The Math:** If a point has 3 samples, the model sees 3 distinct data points. If the lab provides a "Golden Standard" value for that same spot, the model treats the lab value as having nearly zero noise ($\sigma^2_{\text{lab}} \approx 0$) and the sensor values as high noise ($\sigma^2_{\text{sensor}} \gg 0$).
 - **The Result:** The model naturally "snaps" to the Lab values while using the noisy sensors to determine the *shape* and *trend* around those points.
-

2. Improving the Workflow: From "XGBoost-to-Web" to "End-to-End Bayesian"

The Current Problem: By digitizing the XGBoost output into "levels" and taking the midpoint, you are throwing away the **Aleatoric Uncertainty** (the noise the model *knows* it has).

The Proposed Workflow:

1. **Feature Extraction:** Use your XGBoost/ML models to output two things instead of one: a **Mean** and a **Variance** (use a Probabilistic Gradient Boosting or a Dropout-based Neural Net).
 2. **Covariate Noise (The Breakthrough):** If Nitrogen and Phosphorus sensors always spike together when the device is "glitching," the model should learn this **Covariance of Noise**.
 3. **Sensor Fusion:** Instead of predicting nutrients individually, use a **Multi-Output GP**. This allows the model to say: *"I'm not sure about the Nitrogen sensor here, but the Moisture and Conductivity sensors are rock solid, and historically they correlate, so I'll trust the Nitrogen estimate more."*
-

3. Open Source Data for "Priors" (Backfilling)

To give your model a "head start" before the farmer even arrives, integrate these:

- **Sentinel-2 (EO Browser / Copernicus):** Multi-spectral satellite data. **NDVI** (Biomass) and **NDRE** (Chlorophyll) are free, updated every 5 days, and correlate heavily with Nitrogen.
 - **LUCAS Soil Database:** A massive open-source dataset of European soil samples. You can use this to set the "Average" for a specific region/soil type in Germany.
 - **SoilGrids (ISRIC):** Provides global 250m resolution maps of pH, organic carbon, and texture. This is your "Base Layer" prior.
 - **Historical Compatibility:** Since satellite data goes back years, you can "backfill" by looking at how a specific field performed in 2023–2025 to predict its 2026 baseline.
-

4. Dealing with "Massive Data" & Fast Updates

You mentioned the concern about speed. \$10,000\$ points is the limit for "exact" math. To handle millions of sensor pings and do "instant" updates:

- **Inducing Points (Sparse GPs):** You don't calculate the relationship between every grain of sand. The model selects 500 "representative" points across the field. Updates then happen in **milliseconds**.
 - **Kalman Filtering for Time-Series:** For the "noisy sensor" part of the device, use a Kalman Filter on the device/backend to clean the time-series *before* it hits the spatial model.
 - **Online Learning:** Bayesian models are designed for this. When a new sample comes in, you don't retrain the whole field; you just update the **Posterior** using the previous **Prior**.
-

Summary for the "Final" Solution

To build the most accurate system possible:

1. **Stop Digitizing:** Pass the full probability distribution (Mean + Error) from your ML models to the spatial engine.
2. **Multi-Fidelity Input:** Treat Lab data as "High-Fidelity" (Truth) and Sensors as "Low-Fidelity" (Trend).
3. **The "Jury" Effect:** Use Satellite data (Sentinel-2) as the "Grand Prior." If the sensor says "Low" but the satellite sees "Green," the model flags it for a re-sample.
4. **Active UI:** Show the farmer the **Uncertainty Map**. If the device is noisy that day, the map turns "fuzzy," and the app says: *"High interference detected. Please take 2 additional samples at Point B to verify."*

What you haven't thought of: Topography. Water and nutrients flow downhill. Adding a Digital Elevation Model (DEM)—also free from satellite data—as a covariate will explain why certain "jumps" in nutrients happen at the bottom of slopes. This makes your "stepped" functions look logical to the model.

Target Uncertainty & Covariant Noise

This is where GPyTorch excels compared to simple "Kriging":

- **Heteroscedastic Likelihood:** You can explicitly tell the model that the noise is not constant. You can pass the noise estimate for each sample (e.g., $1/\sqrt{n}$) so the model knows which points to trust.

- **Multitask Learning (Multivariate GP):** If you want to model Nitrogen, Phosphorus, and Potassium simultaneously, you use an `IntrinsicCoregionalizationModel` (ICM).
- **Covariant Noise:** It can learn the **noise dependency between features**. If Nitrogen and Moisture sensors usually fail together, the model learns a **Multi-target Covariance Matrix**. It uses the "clean" moisture signal to help "rectify" the noisy nitrogen signal.