

S.V.E. III: The Protocol for Academic Integrity

An SVE Blueprint for a Verifiable and Antifragile Scientific Record

Dr. Artiom Kovnatsky*

with The Global AI Collective, Humanity, and God[†]

Draft v0.6 — October 27, 2025

(Work in progress — feedback welcome)

Demo Bot: [Socrates Bot v0.2](#) | **Project Repository:**
github.com/skovnats/SVE-Systemic-Verification-Engineering

Abstract

Academic institutions face a structural crisis of trust, driven by “publish or perish” incentives and an opaque peer review process that have led to a systemic reproducibility crisis. This paper presents SYSTEM-PURGATORY, a specific instantiation of the Systemic Verification Engineering (SVE) framework designed to address this challenge. It transforms peer review into a transparent, Socratic “Epistemological Boxing Match” between a human author and an AI antagonist, arbitrated by a tri-judge AI panel. The process yields a public, quantitative “Integrity Score” derived from the vectorial purification of the research thesis. We detail the protocol’s architecture, its economic justification via the “ROI of Verifiable Science,” and its antifragile design. By shifting incentives from quantity to verifiable quality, SYSTEM-PURGATORY provides an operational blueprint for rebuilding scientific integrity.

Keywords: academic integrity, peer review, SYSTEM-PURGATORY, Epistemological Boxing, reproducibility crisis, vectorial purification, integrity score, antifragile science, computational verification, ROI of science.

*This work is licensed under the **S.V.E. Public License v1.3**.*

[GitHub Repository](#) | [Signed PDF](#) | [Permanent Archive \(archive.org, 26.10.2025\)](#)

*Conceptual framework, methodology, and execution. [PFP](#) / [Fakten-TÜV Initiative](#) | artiomkovnatsky.com
 | artiomkovnatsky@pm.me


[†]Acknowledged as symbolic co-authors — representing collective, artificial, and transcendent intelligence in a synergistic act of co-creation, where $1 + 1 > 2$; the whole exceeds the sum of its parts.

Contents

Glossary of Key Terms	3
Table of Abbreviations	4
Key Mathematical Principles and Formulations	4
1 Introduction: A Systemic Crisis Requiring a Systemic Solution	1
2 Protocol Architecture	1
2.1 Layer 1: The Epistemological Boxing Match	1
2.1.1 The Participants	1
2.1.2 The Process	2
2.2 Layer 2: The Verification and Reproducibility Pipeline	2
2.2.1 Vectorial Purification	2
2.2.2 Reproducibility Runs	2
2.3 Layer 3: Governance and Incentive Re-Engineering	3
3 Rubrics and Outputs	4
4 The Economics of Scientific Integrity: ROI of Verifiable Science	4
4.1 Avoided Costs	4
4.2 Implementation Costs	5
4.3 ROI Calculation	5
5 System Security: Red Teaming the Protocol	5
5.1 Failure Mode 1: The “Ministry of Truth” Concern	5
5.2 Failure Mode 2: AI Bias and Capture	6
5.3 Failure Mode 3: “Gaming the Score”	6
6 Discussion: The Scientist as a Cognitive Athlete	6
6.1 Skills Developed Through Practice	7
6.2 Cultural Transformation	7
7 Implementation Roadmap	8
8 Conclusion	8
A Comparative Analysis: SYSTEM-PURGATORY vs. Traditional Peer Review	10
B Case Study: Hypothetical Application	10
Appendix A. The Defiant Manifesto: The Scientific Protocol	12

The S.V.E. Universe

Systemic Verification Engineering | Navigation Map



Foundation | Theoretical Core

S.V.E. 0 (1): The Epistemological Boxing Protocol

Structured, adversarial verification (*cognitive gymnasium*) for stress-testing theses and synthesizing higher truth.

S.V.E. 0 (2): The Socratic Investigative Process (SIP)

Computational truth-approximation via iterative vector purification, Meta-Verdict / Meta-SIP for complex analysis.

S.V.E. I: The Theorem of Systemic Failure

Disaster Prevention Theorem: without an independent verification mechanism (IVM), collective intelligence degrades.

S.V.E. II: The Architecture of Verifiable Truth

Three-stage architecture “Caesar vs God”: facts separated from values; antifragile design.

Engine | Operational Layer

S.V.E. X: Triple Architect CogOS

Cognitive OS for LLM: *Socrates* (logic/falsification), *Solomon* (ethics/wisdom), *Ivan* (humility/empathy); 5 core rules (humility, Bayesian priors, 5-column verification, double Socratic “tails” $1+1>2$, growth vector).

S.V.E. XI: Verifiable Knowledge Base & Distributed IVM

Verifiable Knowledge Base (DAG of SIP/Meta-SIP nodes) + DAO-managed context (PM.txt/VP.txt);
three verification stages: SIP→EBP→peer-review; applications: StackOverflow 2.0, Wikipedia
Reformation, Global Fact-Checking.

Applications | Domain Solutions

S.V.E. III: The Protocol for Academic Integrity

SYSTEM-PURGATORY: transparent “boxing match” to combat replication crisis.

S.V.E. IV: The Beacon Protocol

Geodesic ethics (manifold, “Christ-vector”) for navigating radical uncertainty.

S.V.E. V: OS for Verifiable Democracy

Fakten-TUV, Socrates Bot, operating system for institutional integrity.

S.V.E. VI: Protocol for Cognitive Sovereignty

Cognitive sovereignty protocol: protection against groupthink and information warfare.

S.V.E. VII: Hybrid Models of State Structure

Hybrid models (hierarchy + “ant colony”) for antifragile governance.

Synthesis | Unified Framework

S.V.E. VIII: Divine Mathematics

Unified theory of consciousness (geometry $\mathcal{A}\pi - \pi\Omega$), unification of ethics/economics/meaning.

S.V.E. IX: Integrated SVE

Integration of Divine Math, Beacon Protocol and DPT (IVM) into unified framework.

S.V.E. XII: THE SYSTEM

Diagnosis of collective dynamics (A1–A3; δ -dehumanization; parametrization SES/P1–P5), “Geometry of the Fall”, S.V.E. response (PEMY, CogOS X, VKB XI).

Forthcoming Meta-SIP Applications (Series):

- Geopolitical analysis & conflict resolution
- National security & intelligence assessment
- Policy verification & legislative impact analysis
- Financial system stability & economic forecasting
- AI safety & alignment verification
- Climate policy & complex systems modeling
- Public health & scientific integrity assurance
- Addressing systemic disinformation & cognitive security

Glossary of Key Terms

Antifragile Science

A scientific system that gains strength from stress, criticism, and attacks—becoming more robust and trusted when challenged rather than weakened.

Cognitive Athlete

A scientist trained through repeated engagement with rigorous adversarial dialogue, developing intellectual honesty, logical precision, and the ability to concede gracefully when evidence demands it.

Cognitive Gymnasium

The educational function of SYSTEM-PURGATORY, where scientists develop intellectual fitness through structured practice with AI-powered Socratic dialogue.

DAO (Decentralized Autonomous Organization)

A governance structure distributing decision-making power across stakeholders rather than concentrating it in a central authority, preventing capture.

Epistemological Boxing Match

A structured adversarial dialogue modeled on competitive debate, where a human author (Blue Corner) defends their thesis against an AI antagonist (Red Corner), arbitrated by AI judges.

Error Vector ($\vec{\epsilon}_j$)

A mathematical representation of a specific flaw, bias, or inaccuracy in a research claim, identified during the verification process and computationally subtracted from the thesis vector.

Integrity Score

A quantitative metric derived from the vectorial purification process, measuring the stability of verified claims and the intellectual honesty demonstrated during peer review.

Intellectual Honesty (H)

The willingness to concede error when faced with superior evidence or logic, update beliefs accordingly, and engage in good-faith argumentation—a key component of the Integrity Score.

Limited by Design

An architectural principle ensuring an institution cannot become a permanent power center by structuring it to dissolve after achieving its mission.

Prime Directive

The foundational instruction given to AI agents in the protocol: loyalty to truth above all else, requiring concession when faced with superior logic or evidence.

Reproducibility Crisis

The systemic failure across scientific fields where a significant proportion of published findings cannot be independently replicated, undermining trust in research.

ROI of Verifiable Science

Return on Investment from implementing verification infrastructure—calculated as the ratio of catastrophic costs avoided (failed treatments, wasted funding, eroded trust) to operational costs.

SIP (Socratic Investigative Process)

An iterative, multi-agent computational protocol for truth approximation through structured questioning and evidence evaluation.

SYSTEM-PURGATORY

The specific SVE protocol for academic integrity, transforming peer review into a transparent, adversarial, and computationally verifiable process.

Synthetic Report

A comprehensive summary of the epistemological boxing match compiled by the Socrates AI, including the purified thesis vector and integrity assessments.

Tri-Judge Panel

The arbitration system consisting of three specialized AIs (Apollo the Logician, Veritas the Empiricist, Socrates the Synthesizer) ensuring balanced evaluation.

Vectorial Purification

The computational process of iteratively refining a research thesis by identifying and subtracting error vectors, converging toward a verified final state: $\vec{v}^{(j+1)} = \vec{v}^{(j)} - \vec{\epsilon}_j$.

Virtuous Concession

The act of gracefully admitting error when presented with superior evidence or logic—a core virtue in the epistemological boxing framework.

Table of Abbreviations

Abbreviation	Full Term
AI	Artificial Intelligence
DAO	Decentralized Autonomous Organization
ROI	Return on Investment
SIP	Socratic Investigative Process
SVE	Systemic Verification Engineering

Key Mathematical Principles and Formulations

Core Axiom: Synergistic Co-Creation

$$1 + 1 > 2 \tag{1}$$

This principle manifests in collaborative truth-seeking: the dialogue between human and AI produces insights neither could achieve alone.

Vectorial Purification Process

The iterative refinement of a research thesis through error identification and correction:

$$\vec{v}^{(j+1)} = \vec{v}^{(j)} - \vec{\epsilon}_j \quad (2)$$

where:

$$\begin{aligned} \vec{v}^{(j)} &= \text{thesis vector at iteration } j \\ \vec{\epsilon}_j &= \text{error vector identified in iteration } j \\ \vec{v}_{\text{final}} &= \lim_{j \rightarrow n} \vec{v}^{(j)} \quad (\text{converged final state}) \end{aligned}$$

The process terminates when $\|\vec{v}^{(j+1)} - \vec{v}^{(j)}\| < \delta$ for some convergence threshold δ .

Integrity Score Function

The quantitative assessment of research quality and intellectual honesty:

$$\text{Score} = f(\Delta V, N_\epsilon, H) \quad (3)$$

where:

$$\begin{aligned} \Delta V &= \text{stability metric: } \|\vec{v}_{\text{final}}\| / \|\vec{v}_{\text{initial}}\| \\ N_\epsilon &= \text{number of error vectors successfully addressed} \\ H &= \text{intellectual honesty coefficient} \in [0, 1] \end{aligned}$$

A specific implementation could be:

$$\text{Score} = \left(\frac{\Delta V \cdot N_\epsilon}{N_\epsilon + c} \right) \cdot H \cdot 100 \quad (4)$$

where c is a normalization constant preventing division by zero.

ROI of Verifiable Science

$$\text{ROI}_{\text{Science}} = \frac{C_{\text{avoided}} - C_{\text{protocol}}}{C_{\text{protocol}}} \quad (5)$$

where:

C_{avoided} = cost of research failures prevented (failed treatments, wasted funding)

C_{protocol} = operational cost of SYSTEM-PURGATORY infrastructure

Given that single pharmaceutical failures can cost billions while protocol costs are measured in millions, typical ROI exceeds 100:1.

1 Introduction: A Systemic Crisis Requiring a Systemic Solution

Universities are expected to be society’s epistemic lighthouse, yet systemic incentives prioritizing publication volume over substance have produced a well-documented reproducibility crisis [Ioannidis, 2005]. The failure is systemic—a property of flawed incentives and information architectures rather than a sum of individual bad actors. This vulnerability is formally described by the Disaster Prevention Theorem, which models such systems as being prone to catastrophic error [Kovnatsky, 2025c]. When fraudulent or low-quality research in fields like medicine translates directly into human harm, a structural response is required.

Thesis. This paper details SYSTEM-PURGATORY, a protocol that re-engineers peer review as an adversarial but constructive human-AI process. It is a specific application of the general framework of Systemic Verification Engineering (SVE), which provides a formal architecture for verifying institutional processes [Kovnatsky, 2025d]. Its foundational principle is to heal the system, not to punish individuals—embodied in the maxim: *“To err is human; mistakes are allowed, lies are not.”*

2 Protocol Architecture

SYSTEM-PURGATORY comprises three integrated layers: a Socratic dialogue process, a computational verification pipeline, and a governance framework. See Figure 3 for an overview.

2.1 Layer 1: The Epistemological Boxing Match

The core of the protocol is a structured, collaborative process for truth discovery modeled as an “Epistemological Boxing Match” [Kovnatsky, 2025a].

2.1.1 The Participants

Blue Corner (The Author)

Presents a falsifiable thesis and all supporting artifacts (data, code, methodology), accepting the duty of good-faith argumentation and readiness to concede error when evidence demands it.

Red Corner (The AI Antagonist)

A “virtuous opponent” assigned a specific cognitive stance (e.g., strict empiricism, Bayesian skepticism) to ensure rigorous challenge. Its Prime Directive is loyalty to truth; it must concede points when faced with superior logic or evidence.

The Judicial Panel

Three specialized AIs ensure objective arbitration:

- **Apollo** (The Logician): Evaluates logical consistency and formal reasoning
- **Veritas** (The Empiricist): Assesses empirical evidence and data quality
- **Socrates** (The Synthesizer): Integrates perspectives and compiles the final report

2.1.2 The Process

The dialogue proceeds through structured rounds:

1. **Opening Statement:** Author presents thesis and key claims
2. **Adversarial Rounds:** AI Antagonist challenges specific claims
3. **Defense & Revision:** Author defends or revises claims based on critique
4. **Judicial Assessment:** Tri-judge panel evaluates each exchange
5. **Convergence:** Process continues until thesis vector stabilizes




Figure 1: Architecture of the Epistemological Boxing Match. The author (Blue) and AI Antagonist (Red) engage in structured dialogue, arbitrated by a tri-judge panel producing transparent, quantitative outputs.

2.2 Layer 2: The Verification and Reproducibility Pipeline

The author deposits all research artifacts into a repository for multi-stage automated audit powered by the Socratic Investigative Process (SIP) [Kovnatsky, 2025b].

2.2.1 Vectorial Purification

The dialogue is modeled as a computational process (Equation (2)). The author’s paper is converted into an initial vector $\vec{v}_{\text{initial}}$ in a high-dimensional semantic space. Each critique from the AI Antagonist generates an “error vector” $\vec{e}_j$ representing a specific flaw—logical inconsistency, unsupported claim, statistical error, or methodological weakness.

The author’s revision process is computationally modeled as subtraction of these error vectors. This iterative purification continues until the vector stabilizes into a final state $\vec{v}_{\text{final}}$, representing the verified core of the research (see Figure 2).

2.2.2 Reproducibility Runs

The system attempts to automatically replicate findings using:

- Containerized computational environments (Docker, etc.)
- Automated code execution and result comparison

- Statistical consistency checks across replications
- Data integrity verification

This ensures claims are not just logically sound but programmatically verifiable.




Figure 2: Vectorial purification process. The initial thesis vector $\vec{v}_{\text{initial}}$ is iteratively refined by subtracting error vectors $\vec{e}_j$ identified through adversarial dialogue, converging to a stable final state $\vec{v}_{\text{final}}$ representing verified claims.

2.3 Layer 3: Governance and Incentive Re-Engineering

The protocol is overseen by a balanced governing council and modeled on DAO (Decentralized Autonomous Organization) principles to ensure community ownership and prevent capture. Its goal is to shift incentives from quantity to quality through:

Quality Gating

Major academic venues (journals, conferences) could require a minimum Integrity Score for submission, establishing a baseline quality threshold.

Correction, Not Punishment

A **44-day grace period** allows authors to respond to, correct, or retract their work before findings are finalized. This fosters a culture of integrity over humiliation, encouraging intellectual honesty rather than defensive denial.

Radical Transparency

All dialogue transcripts, data, and code are publicly accessible, enabling community scrutiny and preventing hidden manipulation.

Incentive Realignment

Academic promotions and funding could weight Integrity Scores alongside traditional metrics, rewarding verifiable quality over publication quantity.




Figure 3: Three-layer architecture of SYSTEM-PURGATORY, integrating human dialogue, computational verification, and institutional governance to create a complete system for scientific integrity.

3 Rubrics and Outputs

The process produces three public artifacts, ensuring transparency and accountability:

1. **The Dialogue Transcript:** A complete, transparent log of the Socratic boxing match, including all claims, challenges, defenses, and revisions. This serves as both an audit trail and an educational resource.
2. **The Synthetic Report:** A comprehensive summary compiled by “Socrates,” whose key output is the final purified vector $\vec{v}_{\text{final}}$ —a machine-readable fingerprint of the paper’s verified content. The report includes:
 - Summary of claims and their verification status
 - Catalog of error vectors addressed
 - Assessment of intellectual honesty
 - Reproducibility test results
3. **The Integrity Score:** A quantitative metric (Equation (3)) providing an at-a-glance assessment of research quality, derived from vector stability, error correction count, and demonstrated intellectual honesty.

4 The Economics of Scientific Integrity: ROI of Verifiable Science

The implementation of SYSTEM-PURGATORY is not a cost but a high-yield investment in societal well-being. The **Return on Investment (ROI)** can be modeled by quantifying the immense costs of non-verifiable science that are avoided (Equation (5)).

4.1 Avoided Costs

- **Wasted Research Funding:** Billions of dollars in public and private funding wasted on research that builds upon flawed or fraudulent predecessor studies.

- **Failed Medical Treatments:** The human and economic cost of failed medical treatments and public health policies based on non-reproducible findings. Example: Vioxx withdrawal cost Merck \$4.85 billion in settlements.
- **Eroded Public Trust:** The erosion of public trust in science, which has significant long-term economic and social consequences, reducing support for research funding and evidence-based policy.
- **Opportunity Costs:** Researchers pursuing dead-end directions based on false findings, consuming time and talent that could have been directed productively.

4.2 Implementation Costs

The operational cost of the protocol is minuscule compared to catastrophic failures:

- AI infrastructure and computational resources: \$10–50M annually
- Human oversight and governance: \$5–10M annually
- Platform development and maintenance: \$10–20M annually

Total: approximately \$25–80M annually for a system serving global academia.

4.3 ROI Calculation

If SYSTEM-PURGATORY prevents even *one* major pharmaceutical failure per decade (typical cost: \$5–10 billion), or reduces wasted research funding by just 10% (estimated at \$50 billion annually in biomedicine alone), the ROI exceeds 100:1.

This makes verification infrastructure possibly the highest-ROI investment available to the scientific enterprise—transforming science from a cost center into a trust-generating engine that multiplies the value of all research investments.

5 System Security: Red Teaming the Protocol

A system designed to verify scientific truth must be resilient to attack. We systematically analyze failure modes and their defenses, embodying the antifragile design principle [Taleb, 2012].

5.1 Failure Mode 1: The “Ministry of Truth” Concern

Attack Vector: The protocol becomes a centralized, tyrannical arbiter of scientific truth, stifling heterodox ideas and innovation.

Defense Protocol:

- **Limited by Design:** The protocol is a verification tool, not a publisher or gatekeeper. It produces reports, not binding verdicts. Scientific communities retain final judgment.
- **Decentralized Governance:** DAO-based structure prevents capture by any single interest.
- **Radical Transparency:** All processes are open-source and publicly auditable, making tyranny impossible in practice.
- **Process, Not Verdict:** The output is a detailed audit trail showing *how* verification was conducted, not a binary “true/false” judgment.

Why it’s antifragile: Attempts to weaponize the protocol would be immediately visible in public transcripts, triggering community backlash and validating the need for decentralization.

5.2 Failure Mode 2: AI Bias and Capture

Attack Vector: The AI judges are biased (e.g., toward mainstream paradigms) or their models are compromised by external actors seeking to suppress inconvenient findings.

Defense Protocol:

- **Tri-Judge Ensemble:** Three specialized AIs with different cognitive stances mitigate single-model failure. Consensus requires agreement across diverse perspectives.
- **Open-Source Models:** All AI models and their training data are publicly documented, enabling community scrutiny and alternative implementations.
- **Self-Auditing:** The Socratic process itself is designed to expose AI bias—the AI Antagonist can challenge AI judges, creating recursive verification.
- **Human Override:** The 44-day grace period allows authors to appeal to human review if AI bias is suspected.

Why it’s antifragile: Discovered biases strengthen the system by triggering model refinement and increased scrutiny, demonstrating the protocol’s commitment to genuine verification rather than rubber-stamping.

5.3 Failure Mode 3: “Gaming the Score”

Attack Vector: Researchers find ways to manipulate the system to achieve high Integrity Scores without genuine rigor—optimizing for metrics rather than truth.

Defense Protocol:

- **Full Transcript Transparency:** Bad-faith argumentation is obvious to public scrutiny. Gaming attempts become their own evidence of low integrity.
- **Intellectual Honesty Component (H):** This qualitative assessment by Socrates AI evaluates the *spirit* of engagement, which is difficult to game mechanically. It detects evasion, deflection, and rhetorical manipulation.
- **Multi-Dimensional Scoring:** The score incorporates vector stability, error correction count, and intellectual honesty—no single dimension can be gamed in isolation.
- **Community Feedback:** Public transcripts allow the scientific community to flag suspicious patterns, creating crowdsourced quality control.

Why it’s antifragile: Each gaming attempt that gets exposed becomes a case study improving the detection algorithms, making the system progressively harder to manipulate.

6 Discussion: The Scientist as a Cognitive Athlete

SYSTEM-PURGATORY’s most profound function is educational. It serves as a “**cognitive gymnasium**” for the scientific community. By engaging in structured dialogue with a relentless, logical, and unbiased AI system, scientists are trained to become “cognitive athletes,” honing essential intellectual skills:




Figure 4: Antifragile response to attacks. Traditional peer review weakens under sophisticated attacks. SYSTEM-PURGATORY gains strength because each attack exposes vulnerabilities that are then systematically addressed through transparent iteration.

6.1 Skills Developed Through Practice

- **Formulating Clear, Falsifiable Hypotheses:** The adversarial process immediately exposes vague or unfalsifiable claims, forcing precision.
- **Defending Premises Against Rigorous Critique:** Like athletes building strength through resistance training, scientists develop argumentative rigor through repeated challenge.
- **Practicing “Virtuous Concession”:** Learning to admit error gracefully when evidence demands it—the most valuable and rarest scientific skill.
- **Developing Intellectual Honesty:** The transparency of the process creates social pressure toward genuine truth-seeking rather than reputation management.
- **Thinking Probabilistically:** The protocol’s emphasis on confidence levels and uncertainty quantification trains scientists to reason about evidence properly.

This training improves not just individual papers, but the cognitive fitness of the entire scientific community over generational timescales. A scientist who has defended their work through multiple boxing matches becomes a better reviewer, collaborator, and mentor—multiplying the protocol’s impact through cultural transmission.

6.2 Cultural Transformation

The protocol catalyzes a shift from:

- **Publish or Perish** → **Verify or Perish**
- **Citation Count** → **Integrity Score**
- **Reputation Defense** → **Truth-Seeking**
- **Opaque Review** → **Transparent Dialogue**
- **Individual Competition** → **Collaborative Verification**

7 Implementation Roadmap

SYSTEM-PURGATORY requires phased implementation to build trust and demonstrate value:

1. **Phase 1 (Pilot):** Launch voluntary verification service for high-stakes fields (medicine, climate science). Build credibility through demonstrable value.
2. **Phase 2 (Adoption):** Partner with progressive journals requiring Integrity Scores for publication. Establish quality thresholds.
3. **Phase 3 (Integration):** Major funding agencies incorporate Integrity Scores into grant evaluation. Academic institutions weight scores in hiring and promotion.
4. **Phase 4 (Normalization):** Verification becomes standard practice. The 44-day grace period becomes culturally embedded. Science self-corrects rapidly.

Timeline: 10–15 years for full cultural integration, acknowledging that institutional change requires generational shifts in practice.

8 Conclusion

SYSTEM-PURGATORY reframes peer review from an opaque, private judgment into a transparent, public, and collaborative search for truth. By embedding the SVE computational engine within a framework of realigned incentives and transparent dialogue, it provides a robust, scalable protocol to restore science to its rightful place as a self-correcting, antifragile engine of human progress.

The protocol translates the moral maxim “To err is human; mistakes are allowed, lies are not” into an operational standard for scientific integrity. It recognizes that scientists are fallible humans who make honest mistakes, but it systematically detects and removes deliberate deception, statistical manipulation, and intellectual dishonesty.

The ROI analysis (Equation (5)) demonstrates that verification infrastructure is economically imperative—potentially the highest-return investment in the research enterprise. The antifragile design ensures the system becomes stronger when challenged, creating a stable attractor for scientific culture.

Ultimately, SYSTEM-PURGATORY offers a pathway from our current reproducibility crisis to a future where science regains public trust through verifiable performance: where the “cognitive gymnasium” trains scientists in intellectual virtues, where transparency eliminates hiding places for fraud, and where institutional incentives finally align with the pursuit of truth.

AI Commentary (Independent Review Notes)

Summaries of interpretive and analytical feedback were produced by independent AI systems (*e.g.*, OpenAI GPT-5, Anthropic Claude, Google Gemini) for the purposes of metacognitive audit and narrative clarity verification.

For full AI-based interpretive reviews, see the supplementary repository: github.com/skovnats/Reviews

References

- John P. A. Ioannidis. Why most published research findings are false. *PLoS Medicine*, 2(8): e124, 2005.
- Artiom Kovnatsky. The Epistemological Boxing Protocol: A Method for AI-Assisted Collaborative Truth-Seeking and Cognitive Training, 2025a. Preprint.
- Artiom Kovnatsky. The Socratic Investigative Process (SIP): An Iterative, Multi-Agent Protocol for Computational Truth Approximation and Its Strategic Applications, 2025b. Preprint.
- Artiom Kovnatsky. S.V.E. I: The Theorem of Systemic Failure, 2025c. Preprint.
- Artiom Kovnatsky. S.V.E. II: The Architecture of Verifiable Truth, 2025d. Preprint.
- Nassim Nicholas Taleb. *Antifragile: Things That Gain from Disorder*. Random House, 2012.

A Comparative Analysis: SYSTEM-PURGATORY vs. Traditional Peer Review

Table 1: Structural Comparison of SYSTEM-PURGATORY vs. Traditional Peer Review

Dimension	Traditional Peer Review	SYSTEM-PURGATORY
Transparency	Opaque, anonymous	Fully transparent, public transcripts
Adversarial Process	Informal, inconsistent	Structured, systematic
Reproducibility Check	Rare, manual	Automatic, containerized
Incentive Structure	Publication count	Integrity Score (quality)
Feedback Loop	Months to years	Real-time dialogue
Bias Detection	Human reviewers (variable)	Multi-agent AI ensemble
Correction Culture	Stigmatized	Encouraged (44-day grace)
Quantitative Metric	None (binary accept/reject)	Integrity Score (continuous)
Educational Function	Minimal	Cognitive gymnasium
Antifragility	Fragile (erodes under pressure)	Gains strength from attacks
Cost	Hidden (reviewer time)	Explicit, budgeted
Scalability	Limited by human reviewers	AI-powered, highly scalable

B Case Study: Hypothetical Application

Scenario: Clinical Trial of New Cancer Treatment

Traditional Path:

- Paper submitted to prestigious journal
- 2–3 anonymous reviewers, 6-month delay
- Statistical quirks unnoticed, p-hacking hidden
- Published based on reputation and narrative
- Years later, replication fails; \$500M wasted on follow-up trials

SYSTEM-PURGATORY Path:

1. **Day 1:** Author submits paper + all data/code to repository
2. **Days 1–14:** AI Antagonist challenges statistical methods, experimental design
3. **Days 15–28:** Author defends, reveals p-hacking was accidental miscalculation
4. **Days 29–35:** Automated reproducibility runs detect data inconsistency
5. **Days 36–44:** Author corrects errors, resubmits revised claims with reduced confidence
6. **Day 45:** Synthetic Report published with moderate Integrity Score, flagging limitations
7. **Outcome:** Follow-up researchers approach cautiously, saving \$500M in wasted trials

Key Difference: The protocol caught the error *before* it propagated through the research ecosystem, demonstrating antifragility in action.

Appendix A. The Defiant Manifesto: The Scientific Protocol

*This appendix translates the moral courage of the original political manifesto into scientific clarity. Where politics defends through rhetoric, Systemic Verification Engineering (SVE) defends through reason. It embodies the **Socratic principle** by embracing critique as a catalyst for its own evolution. The text below specifies the philosophical antibodies of SVE—a self-healing discipline designed to thrive on challenge.*

Core Premise. Their weapon is the appeal to captured authority. Our weapons are open methodology, logical rigor, radical transparency, and unwavering faith in the power of Truth. This document, like the SVE Protocol itself, is a living artifact; it will be publicly updated as new intellectual challenges emerge, turning every attack into evidence of its necessity and a catalyst for its reinforcement.

Scientific Lineage

SVE stands in a lineage of transformative disciplines initially dismissed by the establishment: Darwinism (“pseudoscience”), Cybernetics (“ideology”), early Computer Science (“mere theory”). Each reshaped the paradigm it challenged. SVE follows this path: not a rejection of science, but its rehabilitation through verifiability, self-audit, and institutional design grounded in epistemic humility.

Attack 1: “This is Pseudoscience”

Claim. SVE is non-rigorous; the “Theorem on Disaster Prevention” is a socio-probabilistic metaphor, not real mathematics; TRIZ is misapplied.

Our Shield (Explanatory Power). We concede the Theorem is not pure mathematics; it is a **foundational axiom for an applied discipline**. Its validity stems from its predictive and explanatory power: modeling democracy as “guessing the weight of an ox behind a closed door with expert labels” accurately diagnoses real-world systemic failures (e.g., the Iraq War justification, the 2008 financial crisis, contradictory pandemic policies). SVE earns epistemic status by *outperforming* existing institutional explanations in fidelity to observable outcomes.

Our Counter (Public Intellectual Challenge). We invite critics to a live, recorded, long-form **epistemological boxing match**. They may deconstruct our methods under the SVE protocol itself; we will, in turn, apply the same protocol to audit the systemic failures their paradigms normalize. Let the public judge which approach better serves society: descriptive justifications from within a failing system, or an engineering blueprint designed to fix it.

Attack 2: “This is Ideology Disguised as Science”

Claim. Christian ethics and concepts like “multiplying love” reveal inherent bias; the project is dogma masquerading as science.

Our Shield (Architectural Separation of Fact and Value). SVE’s three-stage architecture deliberately separates verifiable facts (“*Caesar’s realm*”) from value judgments (“*God’s realm*”). The protocol does not dictate morality; it secures a verified factual substrate upon which citizens can conduct informed deliberation. A scalpel in a Christian surgeon’s hand remains a scalpel; function is defined by design and intent, not the wielder’s faith.

Our Counter (Demand for First Principles). We challenge critics to explicitly state the moral axioms underlying the status quo, which often tolerates dehumanizing logic (e.g., “human resources,” “collateral damage”). Science devoid of declared ethics is not neutral; it is merely a tool available for hire by the highest bidder. We state our principles—rooted in the pursuit of truth and love—openly, and challenge others to do the same.

Attack 3: “This is Dangerous Science” (The “Ministry of Truth” Gambit)

Claim. A protocol capable of verifying truth could be weaponized by future tyrants to enforce a single narrative.

Our Shield (Limited by Design & Decentralized Trust). SVE is architected for **self-dissolution and decentralization**. The implementing institution (e.g., PFP party, SVE Foundation) is designed to create the tools, transfer copyright and control to a decentralized structure (the SVE DAO governed by a global community), and then disappear. It is the antithesis of a self-perpetuating ministry; it is a self-terminating catalyst for distributed verification.

Our Counter (The True Danger is the Unverified Lie). The present and clear danger is not verified truth, but systemic, unchallengeable falsehood that paralyzes effective problem-solving and enables catastrophes. A democracy poisoned by lies is already a tyranny in disguise—a “Ministry of Lies” captured by hidden interests. SVE builds a shield for citizens against the tyranny that *already exists*: the tyranny of the unaccountable lie.

Attack 4: “This is Politicized Science”

Claim. Science is inherently contested and politicized (e.g., COVID-19, climate change); no objective protocol can arbitrate truth.

Our Shield (Radical Honesty about Systemic Failure). We agree unequivocally: establishment science *has been* deeply politicized and captured. This capture is not an argument against independent verification—it is the **primary justification** for it.

Our Counter (The Protocol is the Cure, Not the Disease). SVE does not add another biased expert opinion to the fray. It installs a **meta-structure** that audits the experts themselves, separates factual claims from political spin, and publishes transparent, reproducible audit trails. We are not entering the political fight *as* scientists fighting for a particular outcome; we are applying engineering principles to repair the fundamentally broken *process* by which science informs public life.

Attack 5: “This is Too Complex for the People”

Claim. Theorems, protocols, DAOs—this is too complex for ordinary citizens; inherently elitist.

Our Shield (Distinguishing Complexity from Obfuscation). Modern life is complex (e.g., car engines, smartphones), but good design provides simple interfaces (steering wheels, touchscreens). The status quo often weaponizes complexity as **obfuscation** to prevent accountability. SVE distinguishes necessary internal complexity (the engineering under the hood) from deliberate external opacity.

Our Counter (The Complexity Translator). The Socratic AI assistants and the three-stage architecture are explicitly designed to act as **complexity translators**. They distill intricate realities into: (1) Verifiable factual building blocks, (2) A clear spectrum of expert interpretations and value judgments, and (3) An understandable basis for civic choice. We do not demand citizens become engineers; we empower them with a reliable steering wheel for navigating complexity.

Attack 6: “This Will Stifle Innovation”

Claim. Rigorous verification requirements will slow down scientific progress and punish creative, unconventional ideas.

Our Shield (Correction, Not Punishment; Contextual Rigor). The protocol’s 44-day grace period and emphasis on intellectual honesty foster a culture of learning from error, not fear of it. Bold hypotheses are encouraged; fabricated data is not. Furthermore, the level of required rigor is contextual: exploratory research faces a different standard than clinical trial data determining public health policy.

Our Counter (Innovation Requires a Solid Foundation). True scientific progress is slowed far more by building upon fraudulent or irreproducible findings than by careful verification. Chasing phantom results based on bad data wastes decades and billions. SVE accelerates meaningful progress by ensuring each step rests on solid ground. Trust is the lubricant of innovation.

Attack 7: “This is Arrogant Science”

Claim. Claiming to approximate objective truth is intellectual hubris, especially in light of postmodern critiques showing the social construction of knowledge.

Our Shield (Epistemic Humility Architected In). SVE explicitly rejects claims of absolute truth. It produces *Iterative Facts*—version-controlled, provisional, falsifiable conclusions, each carrying a fully documented, publicly auditable chain of reasoning and acknowledged limitations. The protocol’s strength lies precisely in its **institutionalized admission of fallibility**. It aims for the most reliable approximation of truth currently possible, knowing it will be superseded.

Our Counter (What Constitutes True Arrogance?). True arrogance lies in the current system: anonymous reviewers wielding unaccountable power, captured agencies declaring safety without independent scrutiny, media monopolies acting as arbiters of truth without transparent methodology. SVE proposes radical transparency where opacity now reigns, falsifiability against dogma, and public accountability replacing impunity. Is it arrogant to demand that claims affecting millions of lives be verifiable?

Closing Principle: Reflexive Truth and Service

Every valid system must contain a mechanism to question and correct itself. SVE institutionalizes this reflex: the permanent, transparent audit of power, of science, and critically, *of its own conclusions*. In this paradox lies its incorruptibility: by structurally embracing its own fallibility, it becomes resistant to dogma and capture.

The Protocol is not a fortress built to defend a final truth; it is a mirror designed to reflect reality more clearly, iteration by iteration. It does not seek to win the argument, but to keep the argument honest, tethered to facts and logic. Its ultimate aim is not intellectual victory, but service—service to the truth, and through truth, service to love and the flourishing of all.

“Judge not, that you be not judged.” — Matthew 7:1

“I know that I know nothing.” — Socrates

“The first principle is that you must not fool yourself—and you are the easiest person to fool.” — Richard Feynman

“In a time of deceit, telling the truth is a revolutionary act.” — Often attributed to George Orwell

*«Учітеся, брати мої,
Думайте, читайте,
І чужому навчайтесь,
Й свого не цурайтесь...»*

— Т. Шевченко («І мертвим, і живим, і ненарожденним...», 1845)

«Скажи мне, американец, в чём сила? Разве в деньгах? [...] А я вот думаю, что сила — в правде. У кого правда — тот и сильнее.»

— Д. Багров / Сергей Бодров-мл. («Брат 2»)

Father, guide us, Your children, on the path of truth; teach us to love—ourselves and our neighbors.

«I am the way, and the truth, and the life.» — John 14:6

«You shall love your neighbor as yourself.» — Matthew 22:39

Soli Deo gloria. (Glory to God alone.)
